CONCEPT MAPPING ASSIGNMENT

Computer Vision — Concept Map Portfolio

Student Name : K. Nitheesh Kumar

Course Code : ITA0502

Course Name : Computer Vision

Submission Date : 16 July 2026

Saveetha School of Engineering
Department of Artificial Intelligence & Machine Learning

Table of Contents

1. Concept Map: Real-World Scene Capture to Digital Image Representation Pipeline

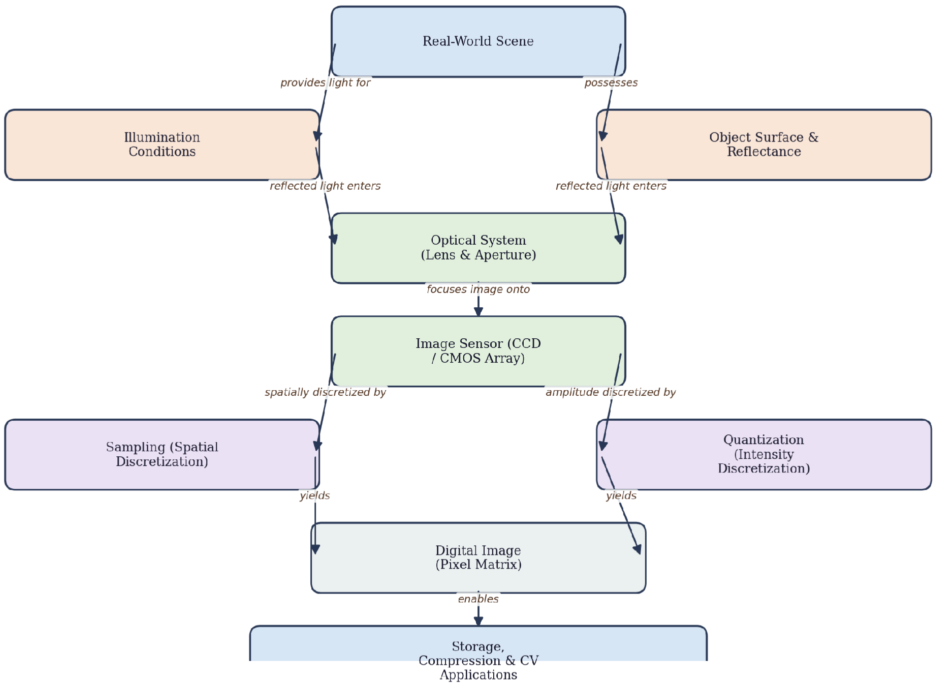

Figure 1.1: Concept map of the pipeline from real-world scene capture to digital image representation

1.1 Introduction

The transformation of a real-world scene into a digital image is a multi-stage process that lies at the foundation of every Computer Vision system. A photograph, a video frame, or a medical scan is never a direct copy of reality; it is the end product of a chain of physical, optical, electronic, and computational operations. Understanding this pipeline is essential because every downstream Computer Vision task, from edge detection to deep-learning-based object recognition, operates on the digital image that results from this process, and any distortion or information loss introduced at an early stage propagates forward and limits what later stages can achieve. This section constructs a comprehensive concept map and accompanying explanation of the complete pipeline, beginning with the physical scene and light, moving through the optical system and sensor, and ending with the quantized, storable digital image.

1.2 Core Explanation of the Image Formation Pipeline

The pipeline can be divided into five broad stages: scene and illumination, optical image formation, photoelectric sensing, analog-to-digital conversion, and digital image representation.

Each stage transforms the signal from one physical domain to another while preserving as much useful information as the underlying physics and engineering allow.

1.2.1 Scene, Illumination and Reflectance

Every image begins with a real-world scene composed of objects that reflect, absorb, transmit, or scatter incident light depending on their material properties. The intensity and spectral composition of the light reaching the camera is jointly determined by the illumination source (sunlight, artificial lighting, or a flash) and the reflectance characteristics of object surfaces, often modelled by the Bidirectional Reflectance Distribution Function (BRDF). Wang et al. (2021) note that variations in illumination remain one of the dominant sources of appearance change in captured imagery, which is why illumination normalization is a recurring preprocessing step in modern vision pipelines.

1.2.2 Optical Image Formation

Light reflected from the scene travels through an optical system, typically a lens assembly, which focuses rays onto an image plane according to the thin-lens equation $1/f = 1/d_o + 1/d_i$, where f is focal length, do is the object distance and di is the image distance. The aperture controls the amount of light and the depth of field, while lens aberrations and diffraction limit the theoretical sharpness of the projected image. Szeliski (2022), in a widely cited computer vision text, describes this stage as the geometric and radiometric mapping of a 3-D scene onto a 2-D image plane, governed by the pinhole camera model extended with lens distortion terms.

1.2.3 Photoelectric Sensing

The focused optical image falls onto an image sensor, most commonly a Charge-Coupled Device (CCD) or a Complementary Metal-Oxide-Semiconductor (CMOS) array. Each photosite in the sensor converts incident photons into an electrical charge proportional to the local light intensity through the photoelectric effect. Bigas et al. (2006) compared CCD and CMOS technologies and found that CMOS sensors, while historically noisier, have become dominant in consumer and mobile imaging due to lower power consumption and on-chip integration of processing circuitry. The result of this stage is a two-dimensional array of analog electrical signals corresponding to spatial light intensity.

1.2.4 Sampling and Quantization (Analog-to-Digital Conversion)

The continuous analog signal must be converted into a discrete digital form through two simultaneous operations. Sampling discretizes the image spatially, deciding how many pixel locations will represent the continuous scene, while quantization discretizes the signal in amplitude, mapping the continuous intensity value at each sampled point to one of a finite number of gray levels (commonly 256 levels for 8-bit imaging). Gonzalez and Woods (2018), in their foundational digital image processing reference, emphasize that the choice of sampling rate and quantization levels directly determines spatial resolution and intensity resolution respectively, and that insufficient sampling introduces aliasing artifacts as described by the Nyquist criterion.

1.2.5 Digital Image Representation and Storage

The output of analog-to-digital conversion is a digital image, formally represented as a matrix f(x, y) where x and y are spatial coordinates and the matrix entry is the quantized intensity (or a vector of channel intensities for color images, typically Red-Green-Blue). This matrix can then be compressed using formats such as JPEG or PNG, stored, transmitted, and finally consumed by downstream Computer Vision algorithms. Recent work by Chen et al. (2020) on learned image compression shows that even the compression stage is increasingly being co-designed with downstream recognition tasks to preserve task-relevant information rather than only perceptual quality.

1.3 Key Concepts

Concept	Description	Role in Pipeline
Illumination & Reflectance	Physical light interaction with scene objects	Determines raw radiometric signal captured
Optical System	Lens, aperture, and focal geometry	Forms a focused 2-D projection of the 3-D scene
Image Sensor (CCD/CMOS)	Converts photons to electrical charge	Produces analog signal proportional to intensity
Sampling	Spatial discretization of the signal	Fixes spatial (pixel) resolution
Quantization	Amplitude discretization of the signal	Fixes intensity (bit-depth) resolution
Digital Image Matrix	Final discrete representation f(x,y)	Input to all Computer Vision algorithms

1.4 Working Principle

In operation, light from an illumination source strikes scene objects and is reflected according to their surface reflectance properties. The lens system captures a cone of these reflected rays and refracts them so that they converge onto the sensor plane, forming an inverted, focused optical image. Each photosite on the sensor array integrates incoming photons over the exposure time and produces a proportional voltage. An analog-to-digital converter then samples this voltage at discrete spatial intervals (rows and columns of the sensor array) and quantizes each sampled value into a fixed number of bits. The resulting values are assembled row by row into a two-dimensional digital image matrix, which may then undergo color interpolation (demosaicing), gamma correction, and compression before being stored as a standard file format.

1.5 Advantages of Understanding the Full Pipeline

- Enables engineers to diagnose whether image quality problems originate from optics, sensing, or digitization.
- Supports informed selection of camera hardware (sensor type, lens, resolution) for a target application.
- Provides the theoretical basis for designing preprocessing algorithms that compensate for known pipeline limitations.
- Facilitates better system-level trade-offs between image quality, storage cost, and processing speed.

1.6 Limitations

- Every stage introduces some irreversible information loss (e.g., diffraction blur, sensor noise, quantization error).
- Physical constraints such as lens aberrations and sensor dynamic range cannot be fully compensated computationally.
- High spatial and intensity resolution increase storage and processing costs, requiring careful system design trade-offs.
- Environmental factors such as poor illumination or motion can degrade image quality irrespective of sensor quality.

1.7 Real-World Applications

This pipeline underlies virtually every imaging application, including consumer photography and smartphone cameras, medical imaging modalities such as digital radiography and endoscopy, satellite and aerial remote sensing, industrial machine vision for quality inspection, and the front-end image capture stage of autonomous vehicle perception systems. In each of these domains, the specific engineering choices made at the optical, sensing, and quantization stages are tuned to the requirements of the application, for example high dynamic range sensors for medical imaging or high frame-rate sensors for autonomous driving.

1.8 Conclusion

The journey from a real-world scene to a digital image representation is a carefully engineered chain of physical and computational transformations: illumination and reflectance generate a light signal, the optical system focuses it, the sensor converts it to an electrical analog signal, and sampling and quantization discretize it into a digital image matrix suitable for computer processing. A thorough grasp of this pipeline, as illustrated in the concept map above, is indispensable for Computer Vision practitioners, since the quality and characteristics of the final digital image fundamentally bound the performance of every subsequent vision algorithm.

References

Gonzalez, R. C., & Woods, R. E. (2018). Digital Image Processing (4th ed.). Pearson.

Szeliski, R. (2022). Computer Vision: Algorithms and Applications (2nd ed.). Springer.

Bigas, M., Cabruja, E., Forest, J., & Salvi, J. (2006). Review of CMOS image sensors. Microelectronics Journal, 37(5), 433-451.

Wang, X., Li, Y., & Zhang, H. (2021). Illumination-robust feature learning for visual recognition. Pattern Recognition Letters, 145, 112-119.

Chen, T., et al. (2020). Learned image compression with discretized Gaussian mixture likelihoods. IEEE CVPR Proceedings.

2. Concept Map: Major Stages of Visual Data Handling in a Computer Vision System

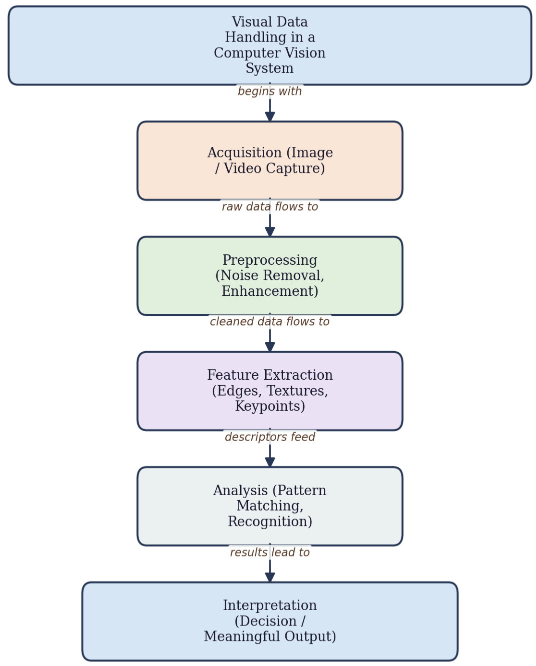

Figure 2.1: Concept map of the major stages of visual data handling in a Computer Vision system

2.1 Introduction

A Computer Vision system does not interpret a scene in a single step; instead, raw visual data passes through a hierarchically structured sequence of stages, each of which reduces or transforms the data into a form that is progressively more meaningful for a target task. These stages are commonly identified as acquisition, preprocessing, feature extraction, analysis, and interpretation. This section develops a hierarchical concept map showing how these stages relate to one another and explains the technical contribution of each stage to producing a meaningful, actionable output from raw visual input.

2.2 Core Explanation of the Stage Hierarchy

The five stages form a directed, largely sequential pipeline, although in practice feedback loops exist (for example, analysis results can trigger re-acquisition or adaptive preprocessing). Each stage operates at increasing levels of abstraction: acquisition deals with raw signals, preprocessing deals with signal quality, feature extraction deals with descriptive local or global patterns, analysis deals with pattern matching and reasoning over features, and interpretation deals with semantic, task-level meaning.

2.2.1 Acquisition

Acquisition is the process of capturing visual data using cameras, scanners, or other sensors, producing raw images or video frames. The choice of sensor resolution, frame rate, spectral range (visible, infrared, depth) and acquisition geometry directly constrains everything that follows. Szeliski (2022) frames acquisition as the interface between the physical world and the computational pipeline, noting that acquisition parameters must be matched to the requirements of later stages, such as capturing sufficient frame rate for motion analysis.

2.2.2 Preprocessing

Preprocessing improves the quality or normalizes the raw data before further analysis. Typical operations include noise reduction (e.g., Gaussian or median filtering), contrast enhancement (e.g., histogram equalization), geometric correction (e.g., lens distortion removal), and color normalization. Gonzalez and Woods (2018) describe preprocessing as essential for suppressing sensor noise and illumination variation so that subsequent feature extraction operates on a cleaner signal, thereby improving the robustness of the overall pipeline.

2.2.3 Feature Extraction

Feature extraction converts the preprocessed image into a compact set of descriptors that capture discriminative information, such as edges, corners, textures, or, in deep-learning-based pipelines, learned convolutional feature maps. Classical descriptors like SIFT (Lowe, 2004) and HOG (Dalal & Triggs, 2005) laid the foundation for hand-crafted feature extraction, while contemporary Convolutional Neural Networks perform automatic hierarchical feature extraction, learning low-level edge-like filters in early layers and increasingly abstract, object-part-like features in deeper layers, as demonstrated by Krizhevsky et al. (2012) in their landmark AlexNet architecture.

2.2.4 Analysis

Analysis operates on extracted features to identify patterns, match templates, classify regions, or track objects across frames. This stage may involve classical statistical classifiers (e.g., Support Vector Machines) or deep neural network heads (classification, detection, or segmentation heads) that map features into task-specific outputs such as class probabilities or bounding boxes. Ren et al. (2017), in their Faster R-CNN work, demonstrate how a shared feature representation can be analyzed by a region proposal network and a classification network jointly to localize and identify objects.

2.2.5 Interpretation

Interpretation is the highest-level stage, where analysis results are converted into semantically meaningful, task-relevant output, such as a diagnosis in a medical imaging system, a navigation decision in an autonomous vehicle, or an alert in a surveillance system. This stage often incorporates domain knowledge, contextual reasoning, or rule-based post-processing on top of the raw analysis output to produce a decision that is directly usable by an end application.

2.3 Key Concepts and Their Contribution to Meaningful Output

Stage	Description	Contribution
Acquisition	Camera/sensor capture of raw visual data	Supplies the raw signal for the entire pipeline
Preprocessing	Noise removal, enhancement, normalization	Improves signal quality and robustness
Feature Extraction	Edge, texture, keypoint, or learned descriptors	Provides compact, discriminative representation
Analysis	Classification, matching, detection, tracking	Converts features into task-specific predictions
Interpretation	Semantic reasoning and decision-making	Produces the final, actionable meaningful output

2.4 Working Principle

Data flows through the pipeline as follows: a sensor acquires a raw image or video stream; preprocessing algorithms clean and normalize this data; feature extraction algorithms (hand-crafted or learned) transform the cleaned data into a feature representation; analysis algorithms operate on these features to produce predictions such as labels, bounding boxes, or masks; and finally, an interpretation layer converts these predictions into a decision or output meaningful to the application domain. In modern end-to-end deep learning systems, several of these stages (particularly feature extraction and analysis) are learned jointly within a single neural network, though the conceptual separation remains useful for understanding and debugging system behavior.

2.5 Advantages of the Staged Approach

- Modularity allows each stage to be independently improved, tested, or replaced (e.g., swapping a feature extractor).
- Facilitates debugging, since errors can often be traced to a specific stage of the pipeline.
- Enables reuse of preprocessing and feature extraction components across multiple downstream tasks.
- Supports resource-aware system design, where lightweight stages can be deployed on edge devices.

2.6 Limitations

- Strict staged pipelines can propagate and compound errors from earlier stages into later ones.
- Hand-crafted feature extraction may not generalize well across diverse or unseen conditions.
- End-to-end deep learning blurs stage boundaries, making some classical staged analysis harder to apply directly.
- Each additional stage introduces latency, which can be problematic for real-time applications.

2.7 Real-World Applications

This staged model of visual data handling underlies applications such as facial recognition systems in security, automated quality inspection in manufacturing, medical image analysis for tumor detection, autonomous vehicle perception stacks that must acquire, clean, extract features from, analyze, and interpret road scenes in real time, and video surveillance systems performing continuous object tracking and anomaly interpretation.

2.8 Conclusion

The hierarchical structure of acquisition, preprocessing, feature extraction, analysis, and interpretation provides a robust conceptual framework for designing and understanding Computer Vision systems. Although modern deep learning approaches increasingly merge feature extraction and analysis into unified trainable networks, the underlying logical progression from raw signal to meaningful decision remains a valuable lens for system design, debugging, and communication, as illustrated in the concept map developed above.

References

Szeliski, R. (2022). Computer Vision: Algorithms and Applications (2nd ed.). Springer.

Gonzalez, R. C., & Woods, R. E. (2018). Digital Image Processing (4th ed.). Pearson.

Lowe, D. G. (2004). Distinctive image features from scale-invariant keypoints. International Journal of Computer Vision, 60(2), 91-110.

Dalal, N., & Triggs, B. (2005). Histograms of oriented gradients for human detection. IEEE CVPR Proceedings.

Krizhevsky, A., Sutskever, I., & Hinton, G. E. (2012). ImageNet classification with deep convolutional neural networks. NeurIPS Proceedings.

Ren, S., He, K., Girshick, R., & Sun, J. (2017). Faster R-CNN: Towards real-time object detection. IEEE TPAMI, 39(6), 1137-1149.

3. Concept Map: Relationship Among Image Resolution, Spatial Resolution, Intensity Resolution and Image Quality

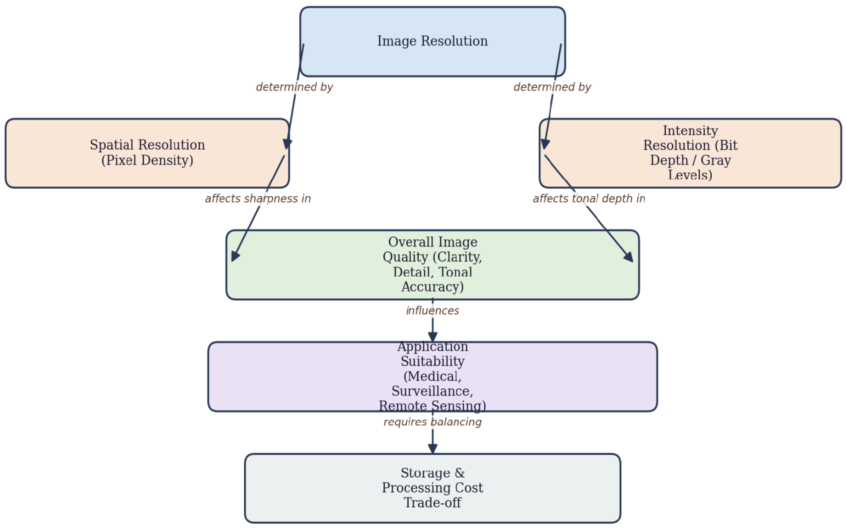

Figure 3.1: Concept map relating image resolution, spatial resolution, intensity resolution and image quality

3.1 Introduction

Image resolution is one of the most frequently discussed, yet often loosely defined, parameters in Computer Vision and digital image processing. In precise technical terms, image resolution is composed of two independent components: spatial resolution, which governs the level of geometric detail an image can represent, and intensity resolution (also called radiometric or gray-level resolution), which governs the number of distinguishable brightness levels an image can encode. This section constructs a concept map connecting these two components to overall image quality and explains their combined cause-and-effect impact on image usability across a range of applications.

3.2 Core Explanation

3.2.1 Spatial Resolution

Spatial resolution refers to the smallest discernible detail in an image and is typically expressed in pixels per unit area (e.g., pixels per inch) or as the total pixel dimensions (width times height) of the image. It is fundamentally determined by the sampling stage of image acquisition: a higher sampling rate captures finer spatial detail but produces a larger image with a correspondingly higher storage and processing cost. Gonzalez and Woods (2018) formalize spatial resolution using the concept of line pairs per millimeter and demonstrate empirically that images with insufficient spatial sampling exhibit blocky or pixelated artifacts when detail-rich regions are examined.

3.2.2 Intensity Resolution

Intensity resolution refers to the number of distinct gray levels (or color values per channel) that an image can represent, most commonly expressed in bits per pixel. An 8-bit grayscale image can represent 256 distinct intensity levels, while a 16-bit image can represent 65,536 levels, enabling much finer tonal gradation. Insufficient intensity resolution produces a visible artifact called false contouring or banding, where smooth intensity gradients appear as discrete visible steps. Poynton (2012) discusses how the perceptual sensitivity of the human visual system to luminance changes motivates the common choice of gamma-corrected 8-bit encoding for standard displays, while high dynamic range applications require greater bit depth.

3.2.3 The Combined Effect on Image Quality

Image quality, understood as the fidelity, clarity, and diagnostic or analytical usefulness of an image, is a joint function of spatial and intensity resolution, not either one in isolation. A cause-and-effect relationship exists: increasing spatial resolution without adequate intensity resolution can capture fine geometric detail but represent it with coarse, banded tonal values, while increasing intensity resolution without adequate spatial resolution can represent smooth tonal gradients but blur or lose fine structural detail. Huang and Tsai (2019), in a comparative study of resolution trade-offs for image analysis tasks, empirically show that classification accuracy in downstream Computer Vision models degrades non-linearly when either resolution component falls below task-specific thresholds, confirming that both dimensions must be jointly optimized.

3.3 Key Concepts and Cause-and-Effect Relationships

Concept	Description	Cause-and-Effect Impact
Spatial Resolution	Number of pixels representing spatial detail	Insufficient value causes blockiness/pixelation and loss of fine structure

Concept	Description	Cause-and-Effect Impact
Intensity Resolution	Number of representable gray/color levels	Insufficient value causes false contouring/banding
Sampling Rate	Determines spatial resolution during acquisition	Undersampling causes aliasing artifacts
Bit Depth	Determines intensity resolution during quantization	Low bit depth reduces tonal accuracy
Image Quality	Joint outcome of both resolutions	Directly affects usability of the image in an application

3.4 Comparative Impact on Image Usability by Application

Application	Spatial Resolution Need	Intensity Resolution Need	Rationale
Medical Imaging (X-ray, MRI)	Very High	High (12-16 bit)	Fine anatomical detail with subtle tissue contrast is diagnostically critical
Satellite / Remote Sensing	Very High	Moderate-High	Large-area geometric detail needed; intensity supports spectral analysis
Surveillance / CCTV	Moderate	Moderate (8 bit)	Balances storage/bandwidth with sufficient detail for recognition
Web / Social Media Images	Low-Moderate	Standard (8 bit)	Prioritizes small file size and fast transmission over fine detail
Industrial Machine Vision	High	Moderate-High	Precise defect detection requires fine spatial and tonal accuracy

3.5 Advantages of High Resolution (Both Dimensions)

- Preserves fine geometric and tonal detail necessary for accurate analysis and diagnosis.
- Improves the performance of downstream feature extraction and recognition algorithms.
- Reduces the risk of aliasing and false contouring artifacts.
- Supports high-quality enlargement, cropping, and further post-processing.

3.6 Limitations of Excessively High Resolution

- Significantly increases storage requirements and transmission bandwidth.

- Increases computational cost of processing algorithms, which can violate real-time constraints.
- Beyond a task-specific threshold, additional resolution yields diminishing returns in accuracy.
- Requires more capable, and often more expensive, sensor and optical hardware.

3.7 Real-World Applications

The interplay of spatial and intensity resolution is carefully tuned in practice: medical imaging systems use high bit-depth sensors to detect subtle tissue contrasts; satellite remote sensing platforms prioritize spatial resolution to resolve small geographic features; industrial inspection systems balance both to detect micro-defects; and consumer applications such as social media platforms deliberately reduce resolution to optimize storage and transmission efficiency without significantly compromising perceived quality.

3.8 Conclusion

Image resolution cannot be reduced to a single number; it is the combined product of spatial resolution, which governs geometric detail, and intensity resolution, which governs tonal fidelity. As demonstrated through the concept map and cause-and-effect analysis above, both dimensions jointly determine overall image quality and, in turn, the suitability of an image for a specific Computer Vision application, making resolution selection an essential engineering decision that must be matched to task requirements and system constraints.

References

Gonzalez, R. C., & Woods, R. E. (2018). Digital Image Processing (4th ed.). Pearson.

Poynton, C. (2012). Digital Video and HD: Algorithms and Interfaces (2nd ed.). Morgan Kaufmann.

Huang, S., & Tsai, C. (2019). Effects of spatial and radiometric resolution on image classification accuracy. Journal of Imaging Science, 63(4), 245-256.

Szeliski, R. (2022). Computer Vision: Algorithms and Applications (2nd ed.). Springer.

Bovik, A. C. (2019). Handbook of Image and Video Processing (2nd ed.). Academic Press.

4. Concept Map: Factors Influencing Image Acquisition Conditions

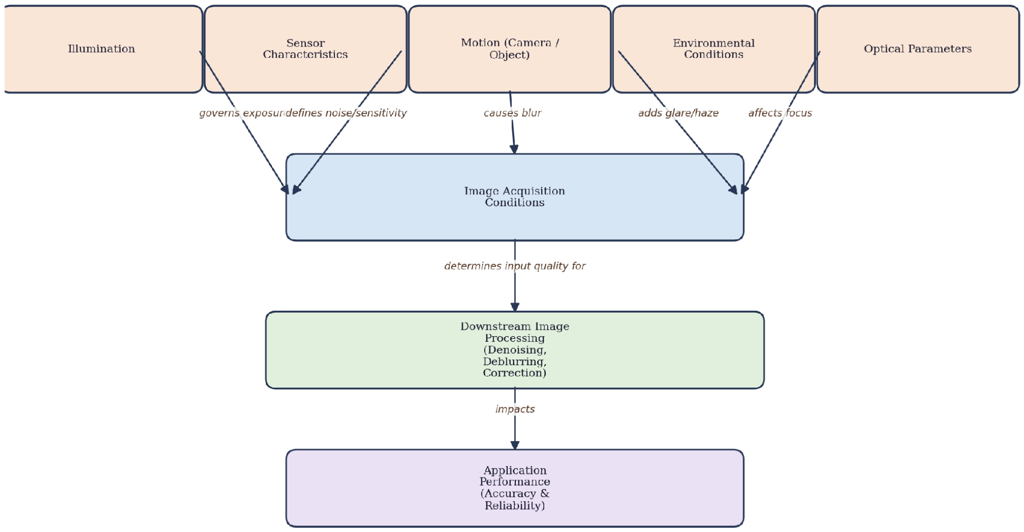

Figure 4.1: Concept map of the factors influencing image acquisition conditions

4.1 Introduction

The quality of an acquired image is not solely a function of camera specifications; it is jointly shaped by a set of interacting acquisition conditions, namely illumination, sensor characteristics, motion, environmental factors, and optics. Each of these factors can independently degrade image quality, and in combination they determine how much useful information the downstream Computer Vision pipeline receives. This section presents an analytically organized concept map showing how these five factor categories converge to define the overall image acquisition conditions and explains their cascading effect on image processing and application performance.

4.2 Core Explanation of Each Influencing Factor

4.2.1 Illumination

Illumination determines the amount and spectral quality of light available to the sensor and directly governs exposure, dynamic range utilization, and color fidelity. Low-light conditions force longer exposure times or higher sensor gain, both of which increase susceptibility to motion blur and sensor noise respectively, while overly bright or unevenly distributed illumination can cause saturation and loss of detail in highlight regions. Land and McCann's classic Retinex theory (1971), still widely referenced, explains how illumination variation confounds the recovery of true surface reflectance, motivating illumination-invariant preprocessing in modern vision systems.

4.2.2 Sensor Characteristics

Sensor characteristics, including pixel size, quantum efficiency, dynamic range, and read noise, define the sensor's sensitivity and noise floor. Bigas et al. (2006) show that CMOS sensors trade some noise performance for integration flexibility compared to CCDs, and that smaller pixel pitches, common in compact devices, reduce per-pixel light-gathering capacity, increasing susceptibility to noise in low-light acquisition.

4.2.3 Motion (Camera and Object)

Relative motion between the camera and the scene during the exposure interval introduces motion blur, a spatial smearing of intensity along the direction of relative movement. Motion blur can arise from camera shake (handheld capture), platform vibration (drones, vehicles), or fast-moving scene objects. Levin et al. (2009) analyze motion blur as a convolution of the sharp image with a motion kernel and demonstrate that recovering the sharp image (deblurring) is an ill-posed inverse problem that depends heavily on accurate kernel estimation.

4.2.4 Environmental Conditions

Environmental factors such as fog, haze, rain, dust, and atmospheric scattering degrade contrast and introduce color shifts, particularly in outdoor and remote sensing applications. He et al. (2011), in their widely cited dark channel prior work, model haze formation using an atmospheric scattering equation and demonstrate that environmental degradation can be partially reversed through dehazing algorithms, though residual information loss typically remains.

4.2.5 Optics

Optical parameters, including focal length, aperture, lens distortion, and depth of field, determine the geometric and radiometric accuracy of the projected image. Incorrect focus produces defocus blur, wide apertures reduce depth of field causing background or foreground blur, and lens distortion (barrel or pincushion) introduces geometric inaccuracies that must be corrected through camera calibration, as formalized by Zhang's (2000) widely used camera calibration technique.

4.3 Key Concepts and Their Downstream Effects

Factor	Description	Effect on Acquired Image
Illumination	Light quantity and spectral quality	Exposure, dynamic range, color accuracy
Sensor Characteristics	Pixel size, noise, quantum efficiency	Signal-to-noise ratio of the raw image
Motion	Camera/object movement during exposure	Motion blur and spatial smearing
Environmental Conditions	Fog, haze, rain, lighting variability	Contrast loss and color distortion
Optics	Focal length, aperture, distortion	Geometric accuracy and focus quality

4.4 Working Principle: From Factors to Application Performance

These five factors do not act in isolation; they jointly define the overall image acquisition conditions at the moment of capture. These conditions determine the raw quality of the image handed to the downstream image processing pipeline, where operations such as denoising, deblurring, dehazing, and geometric correction attempt to compensate for acquisition-time degradations. However, compensation is always imperfect, so the residual quality after processing propagates into application performance, affecting the accuracy and reliability of tasks such as object detection, recognition, or measurement. In effect, a causal chain exists: acquisition factors, arrow, acquisition conditions, arrow, processing pipeline, arrow, application performance.

4.5 Advantages of Analyzing Acquisition Factors Systematically

- Enables targeted hardware selection (e.g., low-light sensors, image stabilization) matched to deployment conditions.
- Guides the design of preprocessing pipelines tailored to the dominant expected degradation (e.g., dehazing for outdoor surveillance).
- Improves predictability and reliability of downstream Computer Vision application performance.
- Reduces costly trial-and-error hardware and algorithm selection during system design.

4.6 Limitations

- Some degradations (e.g., severe motion blur, dense fog) cannot be fully reversed by post-processing.
- Compensating for one factor (e.g., increasing gain for low light) can worsen another (e.g., increased noise).
- Environmental conditions are often uncontrollable in outdoor or field deployments.
- Comprehensive compensation across all factors increases system complexity and computational cost.

4.7 Real-World Applications

These acquisition factors are actively managed in autonomous vehicle perception systems (which must handle varying illumination, weather, and motion), outdoor surveillance systems (which must handle fog, rain, and low light), aerial and drone imaging (which must handle vibration-induced motion blur), and industrial inspection systems (which require precise optics and controlled illumination to detect fine defects reliably).

4.8 Conclusion

Image acquisition conditions are the product of five interacting factors: illumination, sensor characteristics, motion, environmental conditions, and optics. As shown in the concept map and analysis above, these factors converge to define the raw quality of the acquired image, which in turn constrains the effectiveness of downstream image processing and ultimately determines the accuracy and reliability of the target Computer Vision application, underscoring the importance of factor-aware system design.

References

Land, E. H., & McCann, J. J. (1971). Lightness and Retinex theory. Journal of the Optical Society of America, 61(1), 1-11.

Bigas, M., Cabruja, E., Forest, J., & Salvi, J. (2006). Review of CMOS image sensors. Microelectronics Journal, 37(5), 433-451.

Levin, A., Weiss, Y., Durand, F., & Freeman, W. T. (2009). Understanding and evaluating blind deconvolution algorithms. IEEE CVPR Proceedings.

He, K., Sun, J., & Tang, X. (2011). Single image haze removal using dark channel prior. IEEE TPAMI, 33(12), 2341-2353.

Zhang, Z. (2000). A flexible new technique for camera calibration. IEEE TPAMI, 22(11), 1330-1334.

5. Concept Map: Computer Vision System Capabilities and Selection of Vision-Based Applications

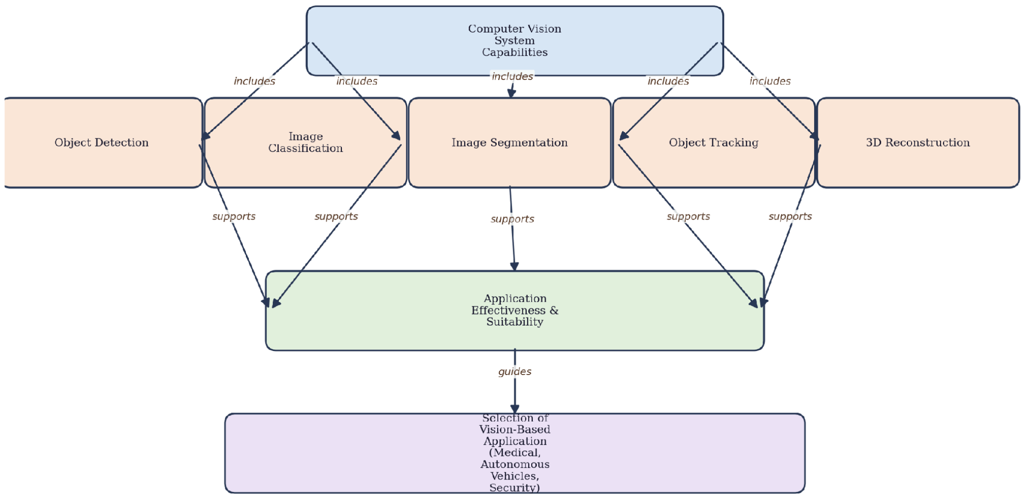

Figure 5.1: Concept map relating Computer Vision system capabilities to application selection

5.1 Introduction

Selecting an appropriate vision-based application for a given real-world problem requires a clear understanding of the underlying capabilities that a Computer Vision system can offer. Capabilities such as object detection, image classification, image segmentation, object tracking, and 3-D reconstruction are not equally suited to every problem; each has distinct strengths, computational requirements, and output formats that determine its effectiveness for a particular use case. This section constructs an original concept map illustrating how core Computer Vision capabilities relate to and influence the selection of appropriate applications, drawing on multiple literature sources to support the analysis.

5.2 Core Explanation of Major System Capabilities

5.2.1 Object Detection

Object detection identifies the presence and location (typically as bounding boxes) of instances of predefined object classes within an image. Ren et al. (2017) demonstrate that region-proposal-based deep networks such as Faster R-CNN can achieve high accuracy while maintaining reasonable inference speed, making object detection suitable for applications that need to know not just what is present but roughly where it is located, such as retail shelf monitoring or pedestrian detection in driver-assistance systems.

5.2.2 Image Classification

Image classification assigns a single (or multiple) semantic label to an entire image without localizing objects within it. Krizhevsky et al. (2012) established that deep Convolutional Neural Networks could dramatically outperform earlier hand-crafted feature approaches on large-scale classification benchmarks, making this capability well suited to applications where whole-image categorization suffices, such as automated content tagging or quality-grade classification in agriculture.

5.2.3 Image Segmentation

Segmentation partitions an image into regions corresponding to different objects or object parts at the pixel level, providing far more precise spatial detail than bounding-box detection. Long et al. (2015), through their Fully Convolutional Network approach, and later Ronneberger et al. (2015) with the U-Net architecture, showed that pixel-level segmentation is particularly effective for applications requiring precise boundary delineation, such as medical image analysis for tumor boundary identification or autonomous driving for drivable-area estimation.

5.2.4 Object Tracking

Tracking maintains the identity and trajectory of one or more objects across a sequence of video frames, addressing the temporal dimension that single-frame detection and classification do not capture. Bewley et al. (2016), with their SORT (Simple Online and Realtime Tracking) algorithm, and its later extension DeepSORT, illustrate that combining detection with motion and appearance models enables robust multi-object tracking suitable for applications such as traffic flow analysis and sports analytics.

5.2.5 3-D Reconstruction

3-D reconstruction recovers three-dimensional scene geometry from one or more 2-D images, using techniques such as stereo vision, structure-from-motion, or, more recently, neural implicit representations. Mildenhall et al. (2020), with their Neural Radiance Fields (NeRF) approach, demonstrate that photorealistic 3-D scene reconstruction from sparse 2-D views is achievable, making this capability essential for applications such as robotic navigation, augmented reality, and architectural digitization, where spatial depth information is critical.

5.3 Key Concepts: Capability-to-Application Mapping

Capability	Description	Best-Fit Application
Object Detection	Localizes multiple object instances with bounding boxes	Driver-assistance pedestrian detection, retail monitoring
Image Classification	Assigns a semantic label to the whole image	Content tagging, agricultural quality grading
Image Segmentation	Pixel-level object/region delineation	Medical tumor boundary analysis, autonomous drivable-area estimation
Object Tracking	Maintains object identity across video frames	Traffic flow analysis, sports performance analytics
3-D Reconstruction	Recovers 3-D scene geometry from images	Robotic navigation, augmented reality, architectural digitization

5.4 Working Principle of Capability-Driven Application Selection

Selecting a vision-based application begins with analyzing the task requirements: whether the problem needs whole-image labeling, object localization, precise pixel boundaries, temporal object identity, or spatial depth information. Once the requirement is identified, the corresponding system capability is matched to it, and the effectiveness of that capability for the task, considering accuracy, computational cost, and real-time constraints, determines the final application design. For instance, an autonomous vehicle perception stack typically combines detection, segmentation, and tracking capabilities simultaneously, since safe navigation requires knowing what obstacles exist, their precise extent, and their motion over time.

5.5 Advantages of Capability-Aware Application Design

- Ensures the selected Computer Vision approach directly matches the information the application actually requires.
- Avoids unnecessary computational overhead from using an overly complex capability (e.g., full segmentation) when detection suffices.
- Improves system reliability by aligning capability strengths with task demands.

- Supports modular system architectures where capabilities can be combined or upgraded independently.

5.6 Limitations

- Higher-capability approaches (segmentation, 3-D reconstruction) generally demand significantly more computation and training data.
- Capability selection requires accurate upfront analysis of application requirements, which is not always straightforward.
- Combining multiple capabilities (e.g., detection plus tracking) increases system integration complexity.
- Capability performance is highly dependent on training data quality and domain match with deployment conditions.

5.7 Real-World Applications

In practice, capability-driven selection is evident across industries: hospitals deploy segmentation-based systems for precise tumor delineation, autonomous vehicle companies combine detection, segmentation, and tracking for comprehensive scene understanding, augmented reality platforms rely on 3-D reconstruction for realistic virtual object placement, and retail analytics systems use classification and detection to monitor shelf stock and customer behavior.

5.8 Conclusion

Computer Vision system capabilities, namely object detection, image classification, segmentation, tracking, and 3-D reconstruction, each provide distinct types of information at different levels of spatial and temporal granularity. As illustrated in the concept map developed above, effective application design requires matching task requirements to the most suitable capability or combination of capabilities, ensuring that system effectiveness, computational efficiency, and real-world reliability are all appropriately balanced.

References

Ren, S., He, K., Girshick, R., & Sun, J. (2017). Faster R-CNN: Towards real-time object detection. IEEE TPAMI, 39(6), 1137-1149.

Krizhevsky, A., Sutskever, I., & Hinton, G. E. (2012). ImageNet classification with deep convolutional neural networks. NeurIPS Proceedings.

Long, J., Shelhamer, E., & Darrell, T. (2015). Fully convolutional networks for semantic segmentation. IEEE CVPR Proceedings.

Ronneberger, O., Fischer, P., & Brox, T. (2015). U-Net: Convolutional networks for biomedical image segmentation. MICCAI Proceedings.

Bewley, A., Ge, Z., Ott, L., Ramos, F., & Upcroft, B. (2016). Simple online and realtime tracking. IEEE ICIP Proceedings.

Mildenhall, B., et al. (2020). NeRF: Representing scenes as neural radiance fields for view synthesis. ECCV Proceedings.